



Observed speed and cost on the same questions



† measured at four concurrent requests per provider, so not comparable with the rows above

The two providers ran at different times through different transports and were not interleaved, so provider load and date are confounders. These describe the observed services on the run dates, not architecture-only speed.